

harinris_HW6_lab

Harin Rishabh

2023-12-10


9.6 Lab: Support Vector Machines

9.6.1 Support Vector Classifier

```
library(e1071)
```

Using the `svm()` function to fit the support vector classifier for a given value of the cost parameter. We begin by generating the observations, which belong to two classes, and checking whether the classes are linearly separable.


```
set.seed(1)
x <- matrix(rnorm(20 * 2), ncol = 2)
y <- c(rep(-1, 10), rep(1, 10))
x[y == 1, ] <- x[y == 1, ] + 1
plot(x, col = (3 - y))
```



We now create a data frame with the response coded as a factor.

```
dat <- data.frame(x = x, y = as.factor(y))
svmfit <- svm(y ~ ., data = dat, kernel = "linear", cost = 10, scale = FALSE)
plot(svmfit, dat)
```

SVM classification plot



We can inspect the support vectors with the index component of the classifier returned from `svm()`.

```
svmfit$index
```


```
## [1] 1 2 5 7 14 16 17
```

```
summary(svmfit)
```

```
##
## Call:
## svm(formula = y ~ ., data = dat, kernel = "linear", cost = 10, scale = FALSE)
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##     cost:  10
##
## Number of Support Vectors:  7
##
## ( 4 3 )
##
##
## Number of Classes:  2
##
## Levels:
## -1 1
```

change the cost parameter and re-run the classifier to see how the support vectors change.

```
svmfit <- svm(y ~ ., data = dat, kernel = "linear", cost = 0.1, scale = FALSE)
plot(svmfit, dat)
```



```
svmfit$index
```

```
## [1] 1 2 3 4 5 7 9 10 12 13 14 15 16 17 18 20
```

We can run cross-validation on a range of values for the cost parameter using the `tune()` function.

```
tune.out <- tune(svm, y ~ ., data = dat, kernel = "linear", ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     5
##
## - best performance: 0.1
##
## - Detailed performance results:
##   cost error dispersion
## 1 1e-03 0.65 0.4743416
## 2 1e-02 0.65 0.4743416
## 3 1e-01 0.15 0.3374743
## 4 1e+00 0.15 0.2415229
## 5 5e+00 0.10 0.2108185
## 6 1e+01 0.10 0.2108185
## 7 1e+02 0.15 0.2415229
```

The best performing model is stored in the `best.model` component of the object returned from `tune()`.

```
bestmod <- tune.out$best.model
summary(bestmod)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = y ~ ., data = dat, ranges = list(cost = c(0.001,
## 0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##         cost: 5
##
## Number of Support Vectors: 8
##
## ( 4 4 )
##
##
## Number of Classes: 2
##
## Levels:
## -1 1
```

We can also classify a different set of observations using the best model obtained from `tune()`. First we create a new set of observations with the `rnorm()` function.

```
xtest <- matrix(rnorm(20 * 2), ncol = 2)
ytest <- sample(c(-1, 1), 20, rep = TRUE)
xtest[ytest == 1, ] <- xtest[ytest == 1, ] + 1
testdat <- data.frame(x = xtest, y = as.factor(ytest))
```

We can then use the `predict()` function to classify the observations in `testdat`.

```
ypred <- predict(bestmod, testdat)
table(predict = ypred, truth = testdat$y)
```

```
##      truth
## predict -1  1
##      -1 10  2
##       1  1  7
```


We can also try a different value for the cost parameter and train the classifier again.

```
svmfit <- svm(y ~ ., data = dat, kernel = "linear", cost = 0.01, scale = FALSE)
ypred <- predict(svmfit, testdat)
table(predict = ypred, truth = testdat$y)
```

```
##      truth
## predict -1  1
##      -1 11  5
##       1  0  4
```

We can linearly separate the two classes and plot the observations.

```
x[y == 1, ] <- x[y == 1, ] + 0.5
plot(x, col = (y + 5)/2, pch = 19)
```




We can also try with a smaller cost value.

```
svmfrit <- svm(y ~ ., data = dat, kernel = "linear", cost = 1)
summary(svmfit)
```

```
##
## Call:
## svm(formula = y ~ ., data = dat, kernel = "linear", cost = 1)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##     cost:  1
##
## Number of Support Vectors:  11
##
##   ( 6 5 )
##
##
## Number of Classes:  2
##
## Levels:
##   -1 1
```


```
plot(svmfit, dat)
```



9.6.2 Support Vector Machine

We begin by generating data with non-linear class boundaries. Plotting the data makes it clear that the class boundary is indeed nonlinear.


```
set.seed(1)
x <- matrix(rnorm(200 * 2), ncol = 2)
x[1:100, ] <- x[1:100, ] + 2
x[101:150, ] <- x[101:150, ] - 2
y <- c(rep(1, 150), rep(2, 50))
dat <- data.frame(x = x, y = as.factor(y))
plot(x, col = y)
```



We split the data into training and test subsets and run the SVM classifier with kernel = “radial” parameter.

```
train <- sample(200, 100)
svmfit <- svm(y ~ ., data = dat[train, ], kernel = "radial", gamma = 1, cost = 1)
plot(svmfit, dat[train, ])
```


SVM classification plot




```
summary(svmfit)
```

```
##
## Call:
## svm(formula = y ~ ., data = dat[train, ], kernel = "radial", gamma = 1,
##     cost = 1)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:    1
##
## Number of Support Vectors: 31
##
## ( 16 15 )
##
##
## Number of Classes: 2
##
## Levels:
## 1 2
```

We can use a larger value for the cost parameter and see if it reduces the training errors.

```
svmfit <- svm(y ~ ., data = dat[train, ], kernel = "radial", gamma = 1, cost = 1e+05)
plot(svmfit, dat[train, ])
```

SVM classification plot



We can run cross-validation using the `tune()` function.

```
tune.out <- tune(svm, y ~ ., data = dat[train, ], kernel = "radial", ranges = list(cost = c(0.1, 1, 10, 100, 1000), gamma = c(0.5, 1, 2, 3, 4)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     1    0.5
##
## - best performance: 0.06
##
## - Detailed performance results:
##   cost gamma error dispersion
## 1 1e-01 0.5 0.27 0.14944341
## 2 1e+00 0.5 0.06 0.08432740
## 3 1e+01 0.5 0.07 0.08232726
## 4 1e+02 0.5 0.10 0.09428090
## 5 1e+03 0.5 0.12 0.10327956
## 6 1e-01 1.0 0.18 0.16865481
## 7 1e+00 1.0 0.07 0.09486833
## 8 1e+01 1.0 0.09 0.09944289
## 9 1e+02 1.0 0.10 0.08164966
## 10 1e+03 1.0 0.09 0.09944289
## 11 1e-01 2.0 0.26 0.15776213
## 12 1e+00 2.0 0.08 0.11352924
## 13 1e+01 2.0 0.12 0.12292726
## 14 1e+02 2.0 0.12 0.13165612
## 15 1e+03 2.0 0.13 0.10593499
## 16 1e-01 3.0 0.27 0.13374935
## 17 1e+00 3.0 0.09 0.11005049
## 18 1e+01 3.0 0.10 0.13333333
## 19 1e+02 3.0 0.13 0.10593499
## 20 1e+03 3.0 0.15 0.10801234
## 21 1e-01 4.0 0.27 0.13374935
## 22 1e+00 4.0 0.09 0.12866839
## 23 1e+01 4.0 0.10 0.13333333
## 24 1e+02 4.0 0.15 0.13540064
## 25 1e+03 4.0 0.17 0.13374935
```

We can predict the classes on the test subset and examine the number of observations misclassified.


```
table(true = dat[-train, "y"], pred = predict(tune.out$best.model, newdata = dat[-train, ]))
```

```
##      pred
## true  1  2
##      1 67 10
##      2  2 21
```

9.6.4 SVM with Multiple Classes

The `svm()` function can also be used to classify observations from multiple-classes.

```
set.seed(1)
x <- rbind(x, matrix(rnorm(50 * 2), ncol = 2))
y <- c(y, rep(0, 50))
x[y == 0, 2] <- x[y == 0, 2] + 2
dat <- data.frame(x = x, y = as.factor(y))
par(mfrow = c(1, 1))
plot(x, col = (y + 1))
```



The `svm()` function now will perform multi-class classification since the dataset we generated now has three class labels.

```
svmfrit <- svm(y ~ ., data = dat, kernel = "radial", cost = 10, gamma = 1)
plot(svmfrit, dat)
```

SVM classification plot

